

# Data Analytics

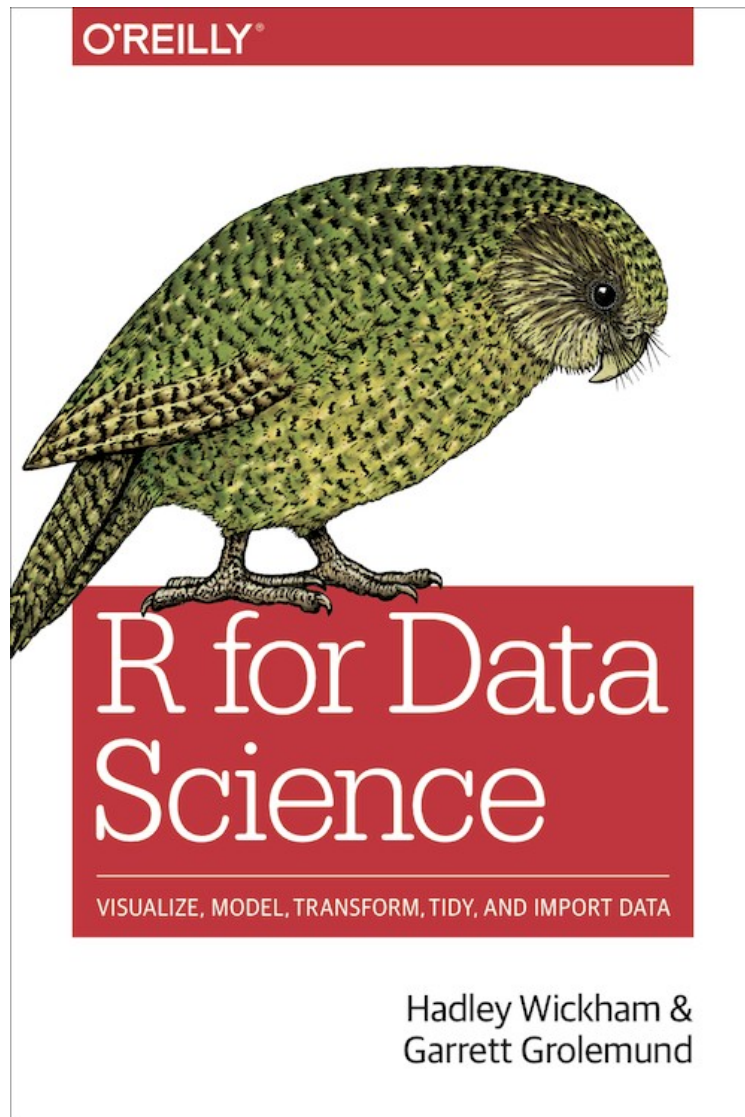
## CS301

### Tidy Data and Import

Week 7: 15<sup>th</sup> Oct  
Fall 2020

Oliver BONHAM-CARTER

# Where in the Web? Where in the Book?

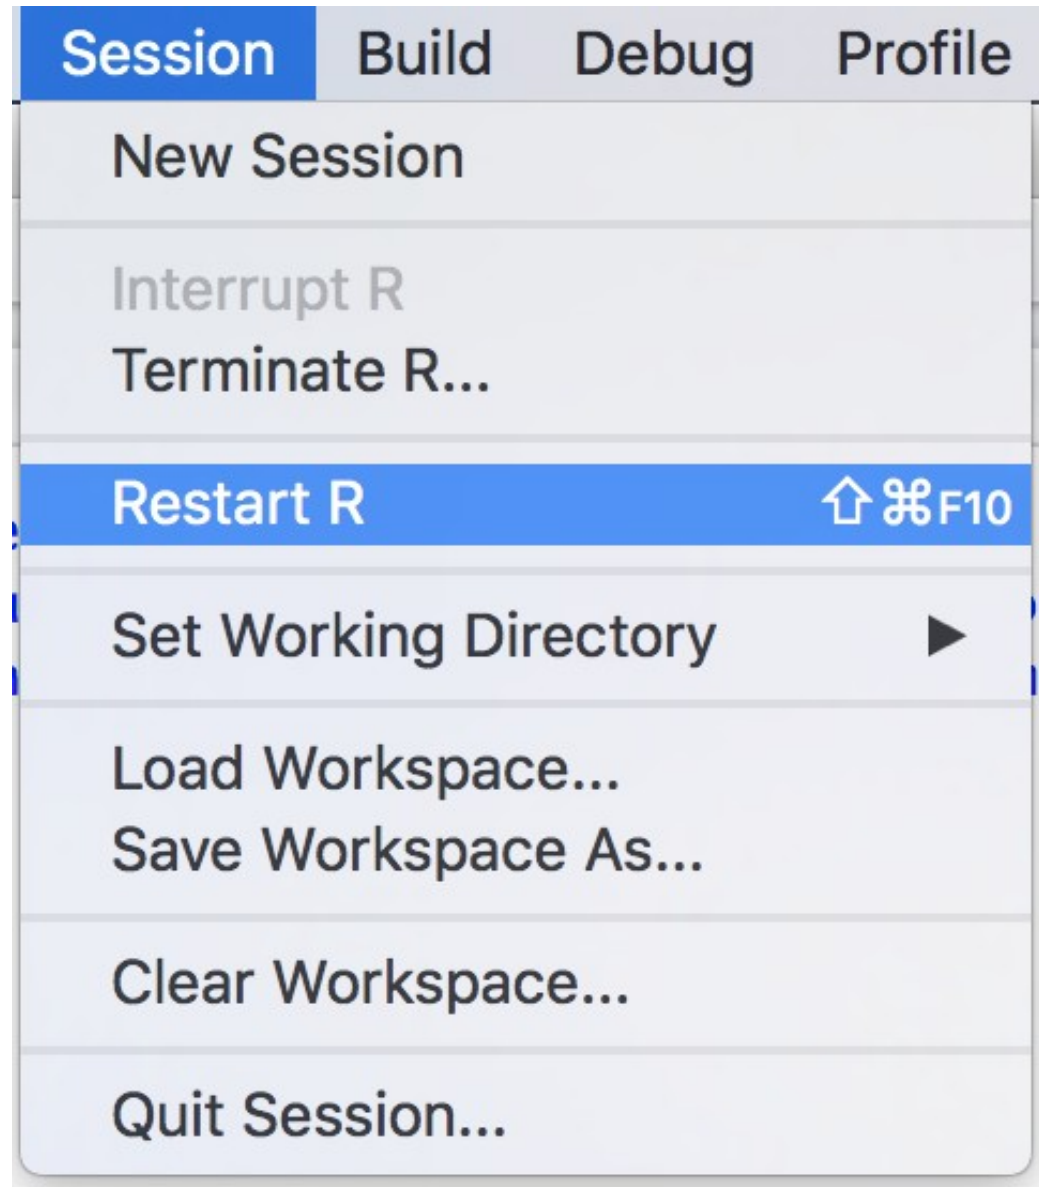


- Note the chapter differences!
- Book:
  - Chap 8
- Web:
  - Chap 11
- Tidy Data and Import



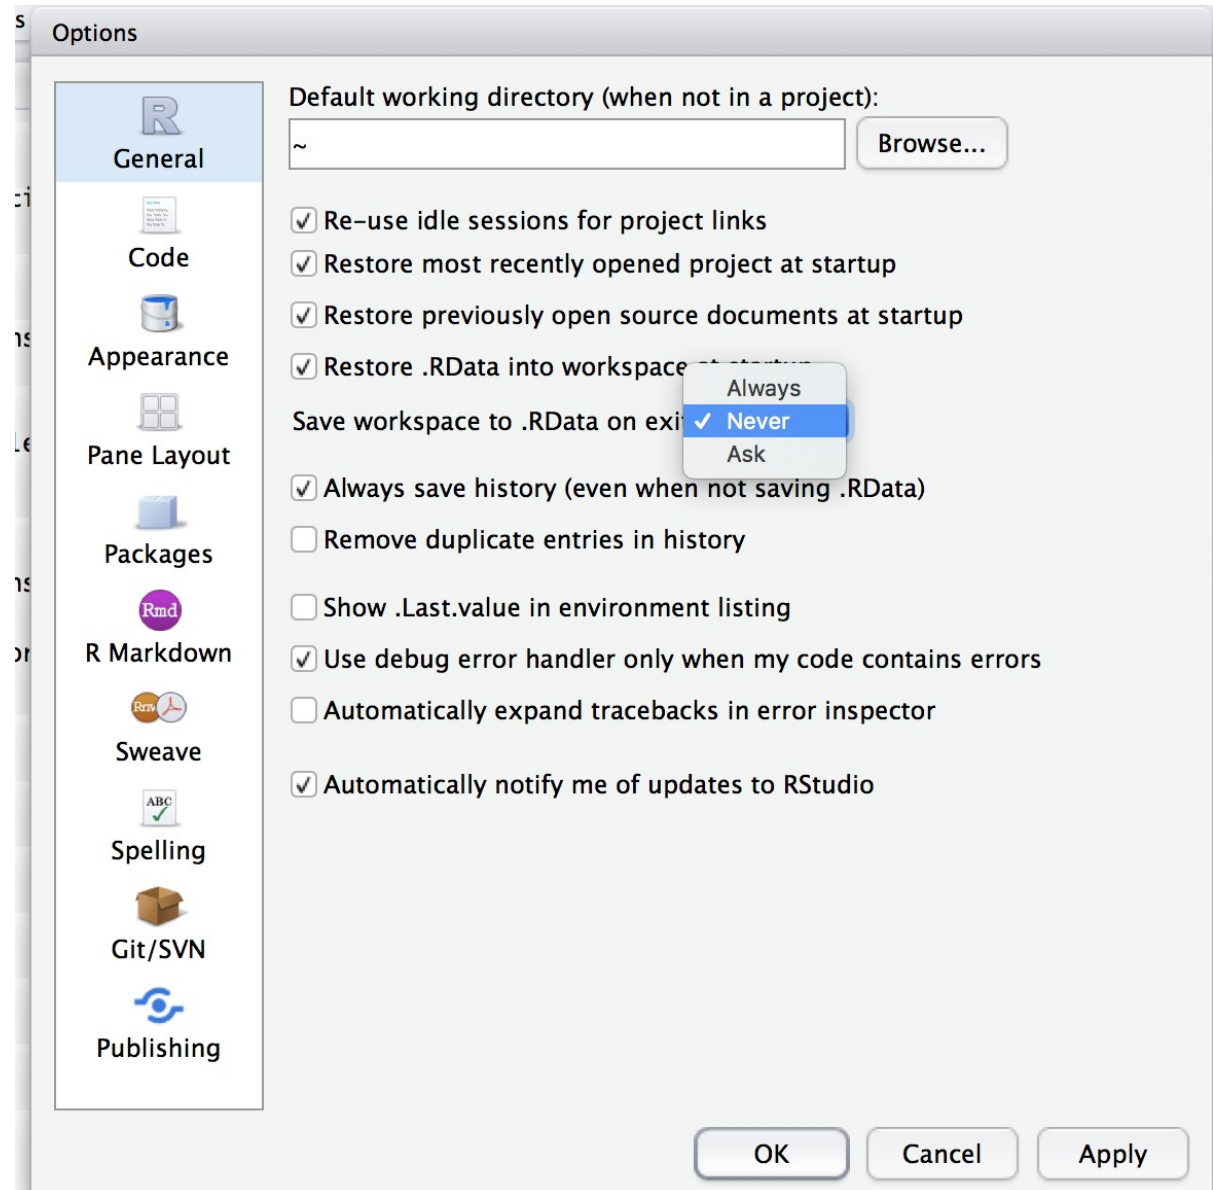
# Now That We Are RStudio Programmers...

- Consider starting with a clean-slate, without a bunch of old data tables.
- Consider not saving your R-environment after each session.
- Instead, your work and code should come source files and not be text-mined from the command history.



# Now That We Are RStudio Programmers...

- Consider stopping the workspace from being saved each time.
- This move will encourage you to begin writing code to be opened in RStudio.
- Better command archive for future works.





# Entering Data as a Table

Your own data typed in:

```
library(tidyverse)
```

```
read_csv("a,b,c  
1,2,3  
4,5,6")
```

Need multiple  
lines to  
define rows

```
read_csv("a,b,c \n 1,2,3 \n 4,5,6")
```

```
read_csv("1,2,3 \n 4,5,6", col_names = FALSE)
```

```
read_csv("1,2,3 \n 4,5,6", col_names = c("col1","col2","col3"))
```

```
> read_csv("a,b,c  
+         1,2,3  
+         4,5,6")  
# A tibble: 2 x 3  
      a     b     c  
  <dbl> <dbl> <dbl>  
1     1     2     3  
2     4     5     6
```

Define column names



# Loading Data and Saving Plots

```
library(tidyverse)

sunSpotData1 <- read.table(file.choose(),
sep=",", header = TRUE)

#sunSpotData2 <- read.table("data/sunSpots.csv",
sep=",", header = TRUE)

#sunSpotData3 <- read_csv("PATH/sunSpots.csv")

ggplot(data = sunSpotData1) + geom_point(mapping =
aes(x = fracOfYear, y = sunspotNum, color = numObs))

#save the plot to file

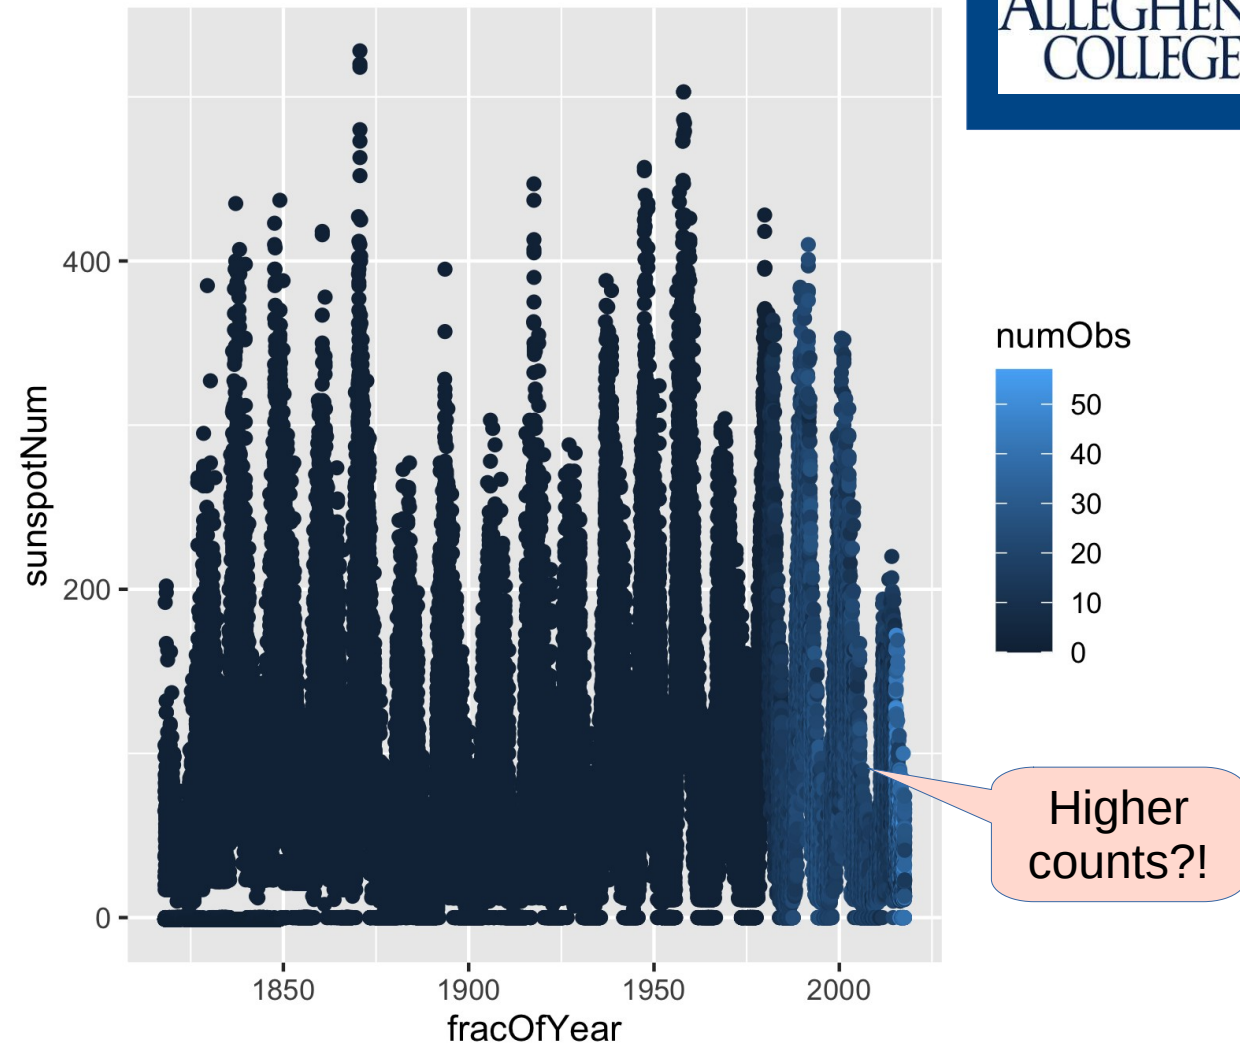
ggsave("~/Desktop/fractOfYearVersusSunspots.png")
```





Wait! Look  
Again!

What does  
this plot tell  
us?

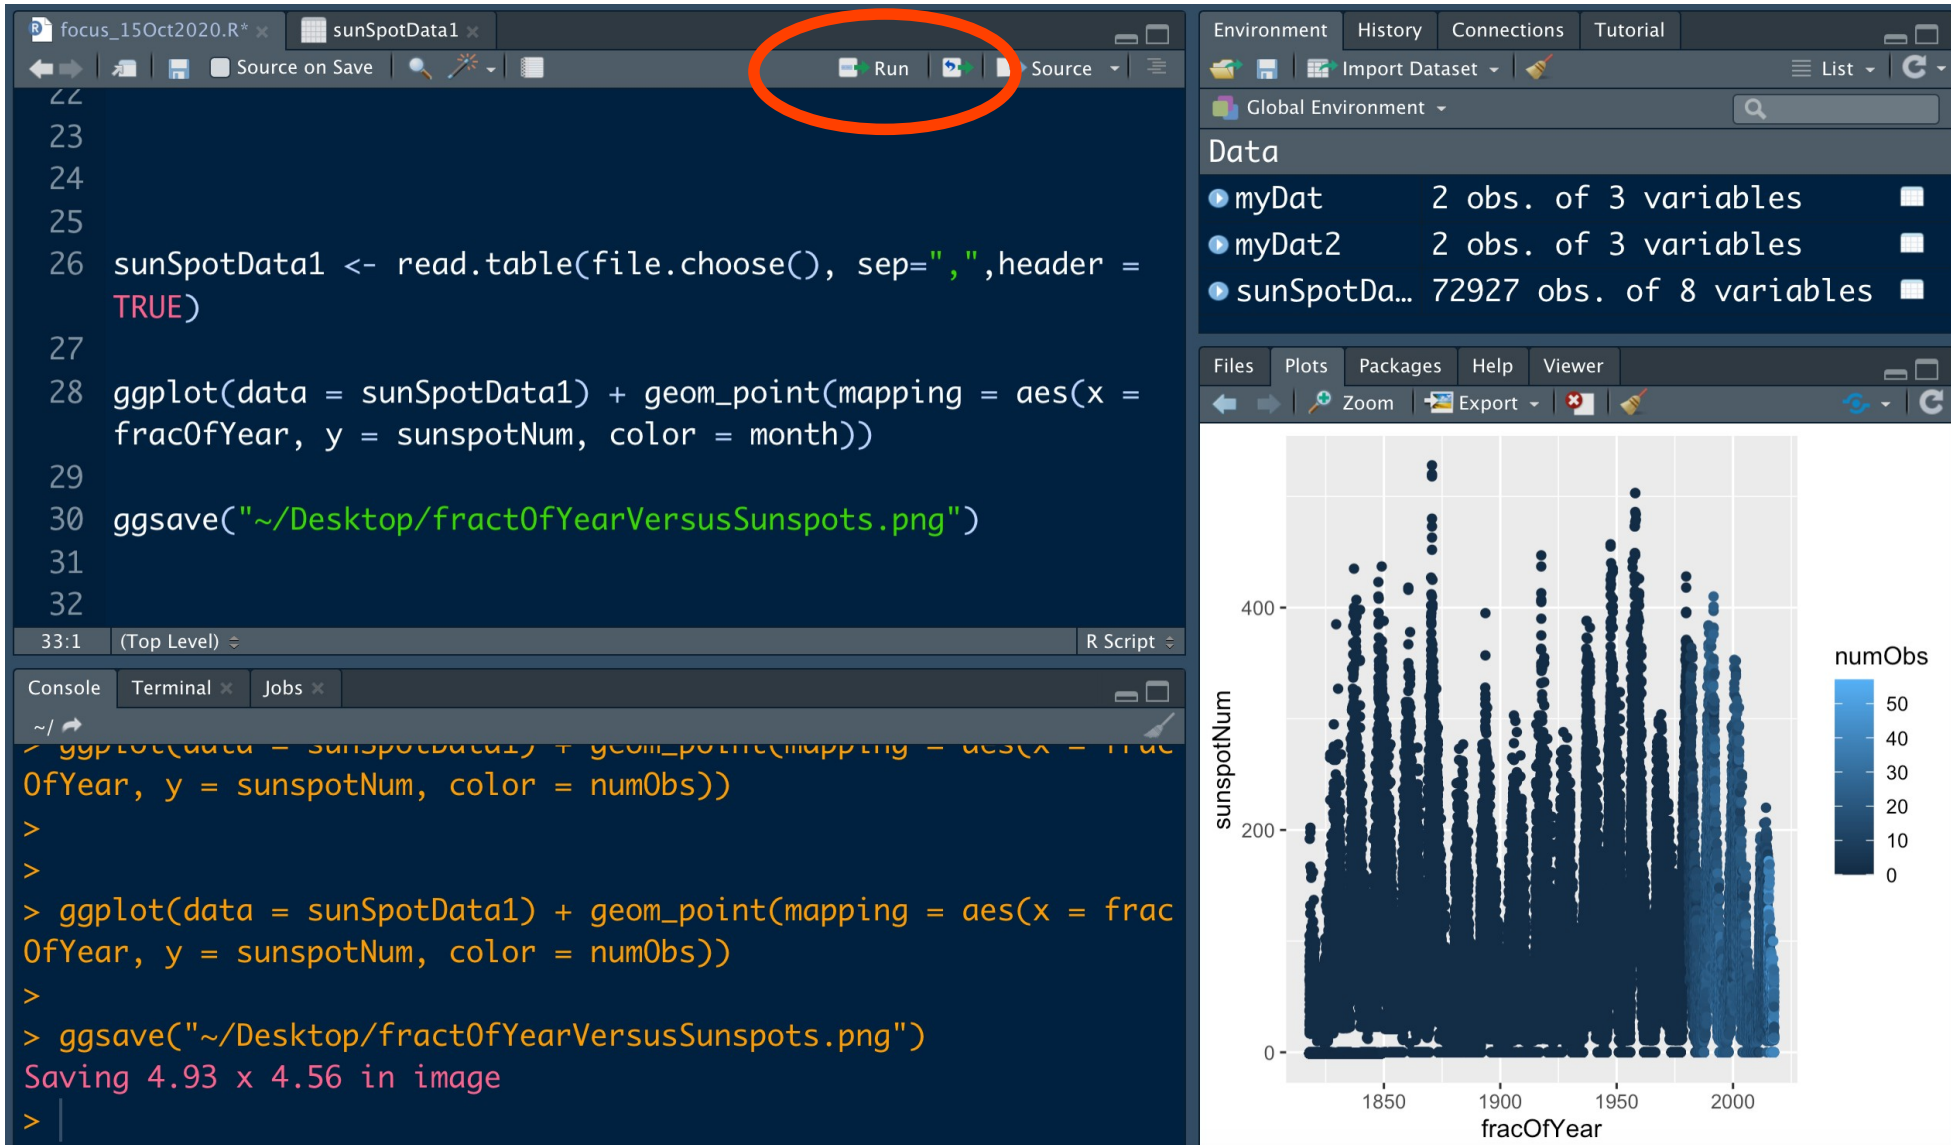


```
ggplot(data = sunSpotData1) + geom_point(mapping = aes(x = fracOfYear,  
y = sunspotNum, color = numObs))
```

```
#save the plot to file
```

```
ggsave("~/Desktop/fractOfYearVersusSunspots.png")
```

Save only good code and  
then have it to run later



The screenshot displays the RStudio environment. The top toolbar features a 'Run' button, which is circled in orange. The main editor window contains the following R code:

```
22  
23  
24  
25  
26 sunSpotData1 <- read.table(file.choose(), sep="," ,header =  
TRUE)  
27  
28 ggplot(data = sunSpotData1) + geom_point(mapping = aes(x =  
fracOfYear, y = sunspotNum, color = month))  
29  
30 ggsave("~/Desktop/fractOfYearVersusSunspots.png")  
31  
32
```

The console window at the bottom shows the execution of the code, with the final command being:

```
> ggsave("~/Desktop/fractOfYearVersusSunspots.png")  
Saving 4.93 x 4.56 in image
```

The right-hand pane shows the 'Environment' tab with a list of objects:

Object	Description
myDat	2 obs. of 3 variables
myDat2	2 obs. of 3 variables
sunSpotDa...	72927 obs. of 8 variables

Below the environment pane, the 'Plots' tab displays a scatter plot titled 'fractOfYearVersusSunspots.png'. The x-axis is labeled 'fracOfYear' and ranges from 1850 to 2000. The y-axis is labeled 'sunspotNum' and ranges from 0 to 400. The plot shows a dense collection of points, with a color scale on the right indicating the number of observations ('numObs') per year, ranging from 0 to 50. The points are colored in shades of blue, with darker blue representing higher values of 'numObs'.



# How Do We Deal With Messy Data?

- We may try to use a data table only to find:
  - There are numbers mixed with characters
  - Different types of entries are mixed in a column
  - Mixed makes things messy.





# The Organization of Data

#Naturally tidy data:

```
library(tibble)
```

```
tibble(x = 1:5, y = 1, z = x ^ 2 + y)
```

What are the qualities  
that make data tidy?!

```
library(tidyverse)
```

# The same data displayed in multiple ways; each data set below organizes the values in a different way

```
table1 # country year cases population
```

```
table2 # country year type count
```

```
table3 # country year rate
```

```
table4a # country `1999` `2000`
```

```
table4b # country `1999` `2000`
```



# Tidy Data

- What does tidy data look like?
  - A column should be of all same types and of same description
- There are three inter-related rules which make a data set *tidy*:
  - Each variable must have its own column.
  - Each observation must have its own row.
  - Each value must have its own cell.

# Tidy Data

- Be tidy: it matters how your data is arranged
- *Trends could be missed due to messy tables*
- Code is easiest to implement when data from a column is same

Figure 9-1 shows the rules visually.

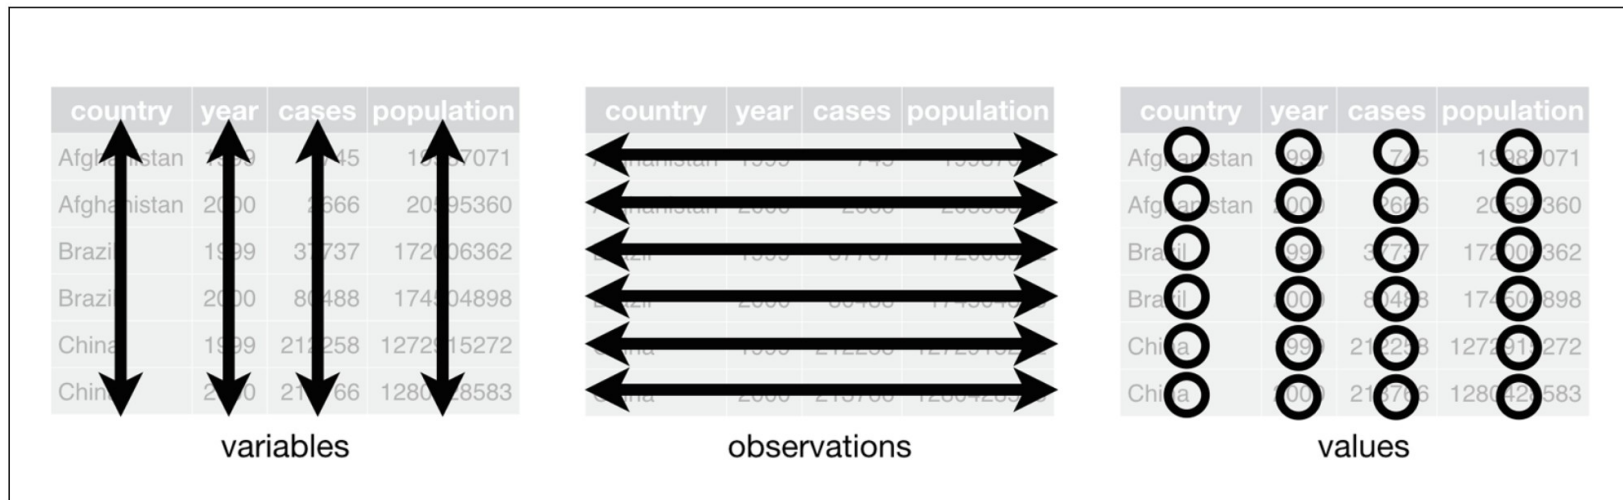


Figure 9-1. The following three rules make a dataset tidy: variables are in columns, observations are in rows, and values are in cells



# Which Table is Most Tidy?

View(table1)

- There are three interrelated rules which make a data set tidy:
  - Does each variable have own column?
  - Does each observation have own row?
  - Does each value have own cell?
- Table 1 is the most tidy for for data-organization

```
> table1
```

```
# A tibble: 6 x 4
```

	country	year	cases	population
	<chr>	<int>	<int>	<int>
1	Afghanistan	1999	745	19987071
2	Afghanistan	2000	2666	20595360
3	Brazil	1999	37737	172006362
4	Brazil	2000	80488	174504898
5	China	1999	212258	1272915272
6	China	2000	213766	1280428583

All same types and descriptions in columns, but it seems that two sets are **mixed**

# Not Tidy!!

- View(table2)
- Not tidy
- The Cases are easily confused



```
> table2
```

```
# A tibble: 12 x 4
```

	country	year	type	count
	<chr>	<int>	<chr>	<int>
1	Afghanistan	1999	cases	745
2	Afghanistan	1999	population	19987071
3	Afghanistan	2000	cases	2666
4	Afghanistan	2000	population	20595360
5	Brazil	1999	cases	37737
6	Brazil	1999	population	172006362
7	Brazil	2000	cases	80488
8	Brazil	2000	population	174504898
9	China	1999	cases	212258
10	China	1999	population	1272915272
11	China	2000	cases	213766
12	China	2000	population	1280428583





Q: What can we do with this data?

A: Count types of observations

- **#Quick Computations** of cases per year (note: wt is the sum of each group of *year*)

```
table1 %>% count(year, wt = cases)
```

```
# <int> <int>
```

```
# 1 1999 250740
```

```
# 2 2000 296920
```

**How many cases for  
1999 and 2000?**

**1999:**  $745 + 37737 + 212258 = 250740$

**2000:**  $213766 + 80488 + 2666 = 296920$

```
# count the populations, aggregated by country
```

```
table1 %>% count(country, wt =  
as.numeric(population))
```



# Implement Ggplots

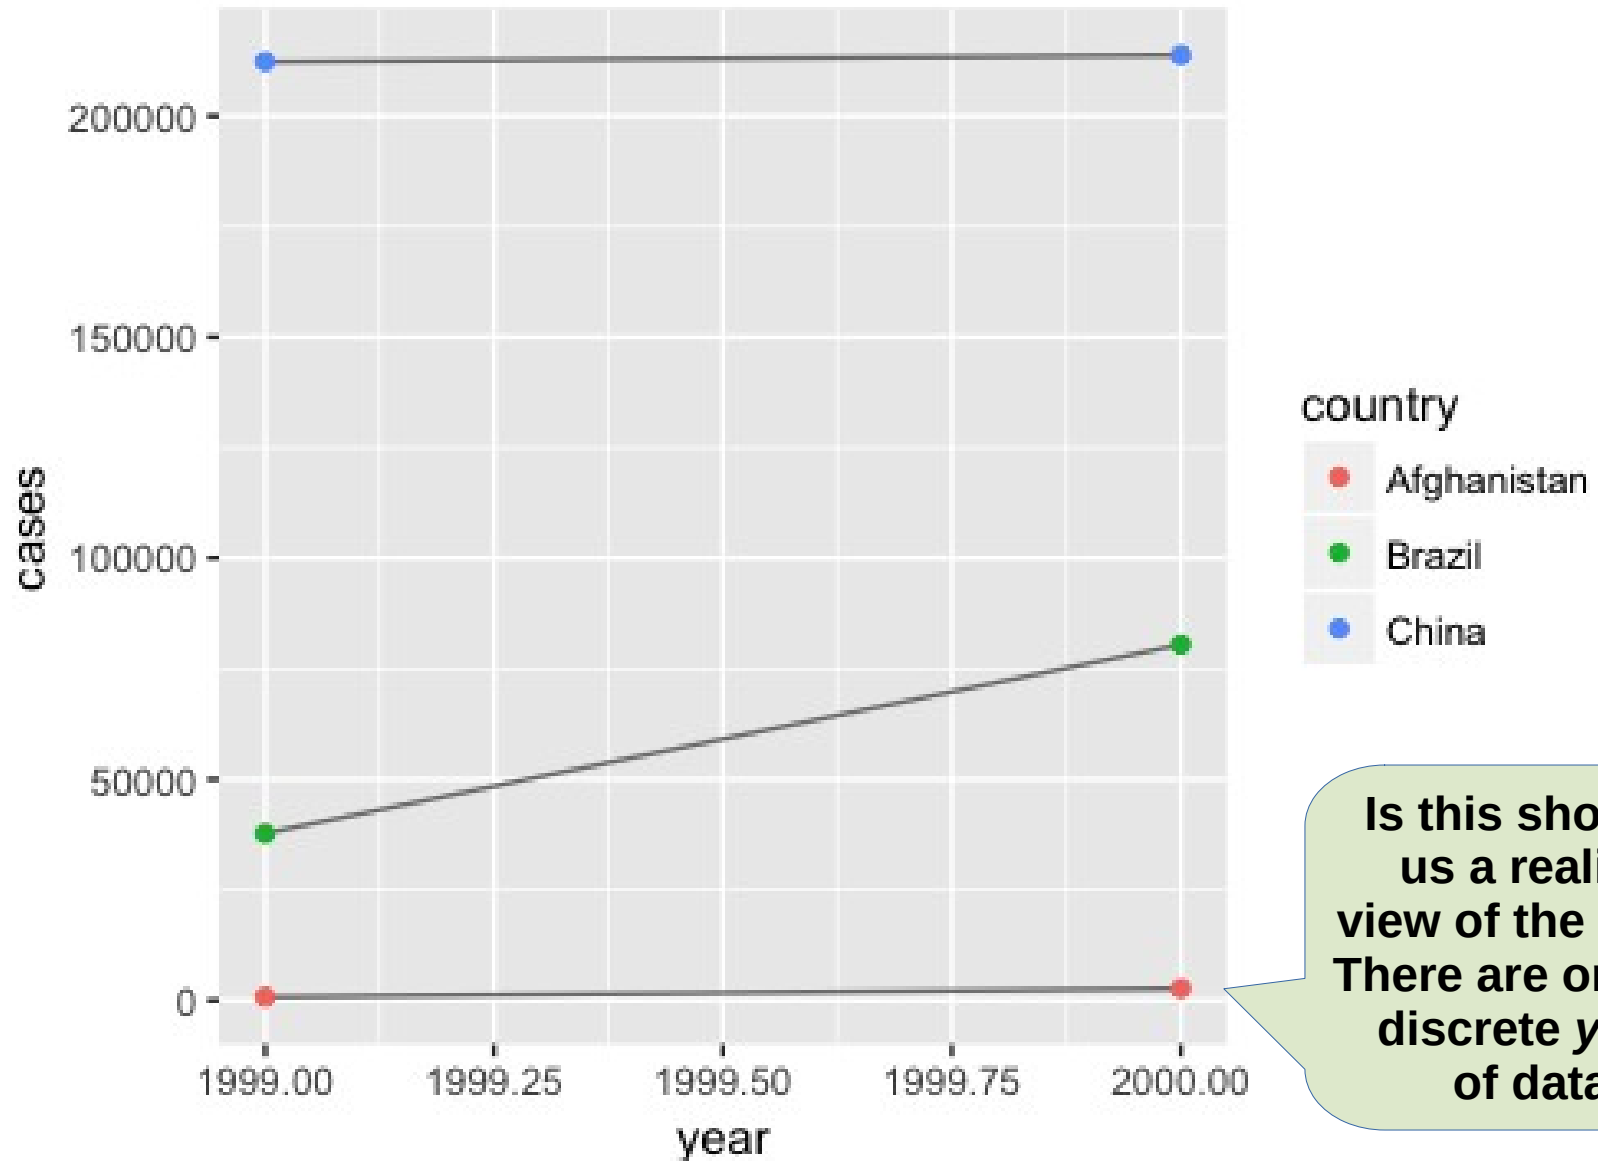
# Visualize changes over time on *table1*

```
ggplot(table1, aes(year, cases)) +  
  geom_line(aes(group = country), colour =  
    "grey50") + geom_point(aes(colour = country))
```

We can still “work” with  
untidy data, right?



# *Discrete Years Become Continuous Years*



Is this showing  
us a realistic  
view of the years?  
There are only two  
discrete years  
of data.

# Bad Organization, Bad Luck!!

- We can apply code to data when in the right format (integers, strings, etc.)
- What happens when the data is badly stored; messy, and without any organization??!





# Gather(): Table4a

- The *gather()* function takes multiple columns to collapse into key-value pairs, duplicating all other columns as needed.
- Use *gather()* when you notice that you have columns that are not variables.

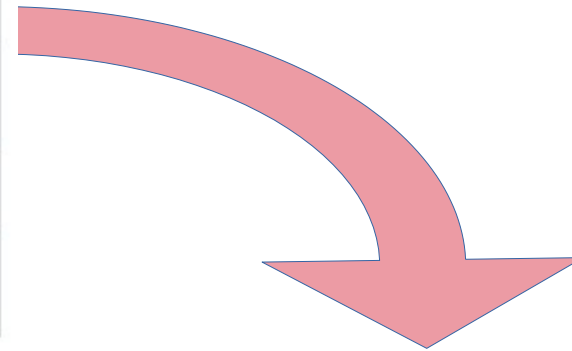
These variables could be better ordered as elements of "Year"

```
> table4a
# A tibble: 3 x 3
  country `1999` `2000`
*      <chr>   <int>   <int>
1 Afghanistan    745    2666
2      Brazil  37737   80488
3      China 212258  213766
```



# All cases data into one column in **table4a**

	country	1999	2000
1	Afghanistan	745	2666
2	Brazil	37737	80488
3	China	212258	213766



```
newTable4a <-  
  table4a %>%  
    gather(`1999`, `2000`,  
          key = "year",  
          value = "cases")
```

```
> table4a %>% gather(`1999`, `2000`,  
key = `year`, value = `cases`)
```

```
# A tibble: 6 x 3
```

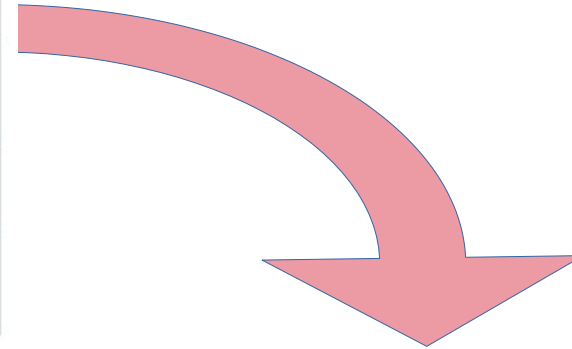
	country	year	cases
	<chr>	<chr>	<int>
1	Afghanistan	1999	745
2	Brazil	1999	37737
3	China	1999	212258
4	Afghanistan	2000	2666
5	Brazil	2000	80488
6	China	2000	213766





# All cases data into one column in **table4a**

	country	1999	2000
1	Afghanistan	745	2666
2	Brazil	37737	80488
3	China	212258	213766



Take the data in columns “1999”  
and “2000” of table4a,  
and place the year into  
new column called “year”.

Then, take the cell data from  
“1999” and “2000” and  
Place it in a new  
column called, “cases”

```
> table4a %>% gather(`1999`, `2000`,  
key = `year`, value = `cases`)
```

```
# A tibble: 6 x 3
```

	country	year	cases
	<chr>	<chr>	<int>
1	Afghanistan	1999	745
2	Brazil	1999	37737
3	China	1999	212258
4	Afghanistan	2000	2666
5	Brazil	2000	80488
6	China	2000	213766



# How did we do that?

```
newTable4a <- table4a %>%  
  gather(`1999`, `2000`,  
    key = "year", value =  
    "cases")
```

Here's how:  
Reorganize the data  
in the columns

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

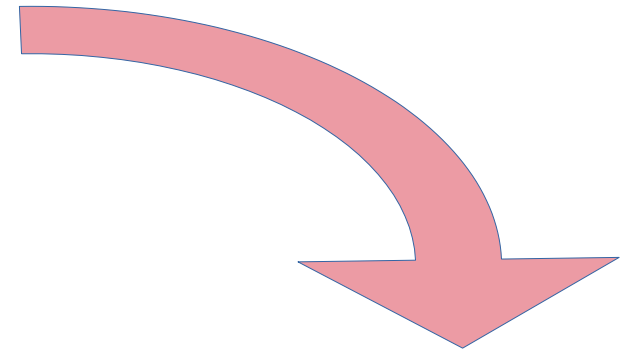
table4

Figure 12.2: Gathering `table4` into a tidy form.



All *population data* into one column in **table4b**

	country	1999	2000
1	Afghanistan	19987071	20595360
2	Brazil	172006362	174504898
3	China	1272915272	1280428583



```
newTable4b <-
```

```
table4b %>%
```

```
gather(`1999`, `2000`,  
  key = "year",  
  value = "population")
```

```
> table4b %>% gather(`1999`, `2000`,  
  key = `year`, value = `population`)
```

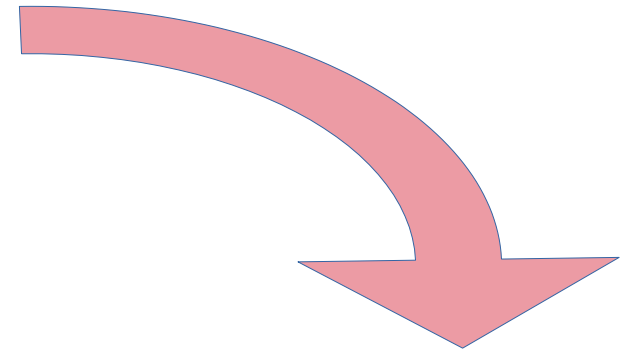
```
# A tibble: 6 x 3
```

	country	year	population
	<chr>	<chr>	<int>
1	Afghanistan	1999	19987071
2	Brazil	1999	172006362
3	China	1999	1272915272
4	Afghanistan	2000	20595360
5	Brazil	2000	174504898
6	China	2000	1280428583



# All *population data* into one column in **table4b**

	country	1999	2000
1	Afghanistan	19987071	20595360
2	Brazil	172006362	174504898
3	China	1272915272	1280428583



Take the data in columns “1999” and “2000” of table4b, and place the year into new column called “year”.

Then, take the cell data from “1999” and “2000” and Place it in a new column called, “population”

```
> table4b %>% gather(`1999`, `2000`,  
  key = `year`, value = `population`)
```

```
# A tibble: 6 x 3
```

	country	year	population
	<chr>	<chr>	<int>
1	Afghanistan	1999	19987071
2	Brazil	1999	172006362
3	China	1999	1272915272
4	Afghanistan	2000	20595360
5	Brazil	2000	174504898
6	China	2000	1280428583





# spread(): table2

- Dealing with mixed values in the same column

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

table2

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

Here's how:  
Reorganize the data  
Into two columns



# spread(): table2

	country	year	type	count
1	Afghanistan	1999	cases	745
2	Afghanistan	1999	population	19987071
3	Afghanistan	2000	cases	2666
4	Afghanistan	2000	population	20595360
5	Brazil	1999	cases	37737
6	Brazil	1999	population	172006362
7	Brazil	2000	cases	80488

Showing 1 to 8 of 12 entries

```
spread(table2, key =  
type, value = count)
```

```
> spread(table2, key = type, value = count)
```

```
# A tibble: 6 x 4
```

	country	year	cases	population
	<chr>	<int>	<int>	<int>
1	Afghanistan	1999	745	19987071
2	Afghanistan	2000	2666	20595360
3	Brazil	1999	37737	172006362
4	Brazil	2000	80488	174504898
5	China	1999	212258	1272915272
6	China	2000	213766	1280428583





# spread(): table2

	country	year	type	count
1	Afghanistan	1999	cases	745
2	Afghanistan	1999	population	19987071
3	Afghanistan	2000	cases	2666
4	Afghanistan	2000	population	20595360
5	Brazil	1999	cases	37737
6	Brazil	1999	population	172006362
7	Brazil	2000	cases	80488

ng 1 to 8 of 12 entries

Take the data in columns “type” and “count” of table2 and place cases data into a single column called “cases” and place the *population* data into new column called “population.”

```
> spread(table2, key = type, value = count)
```

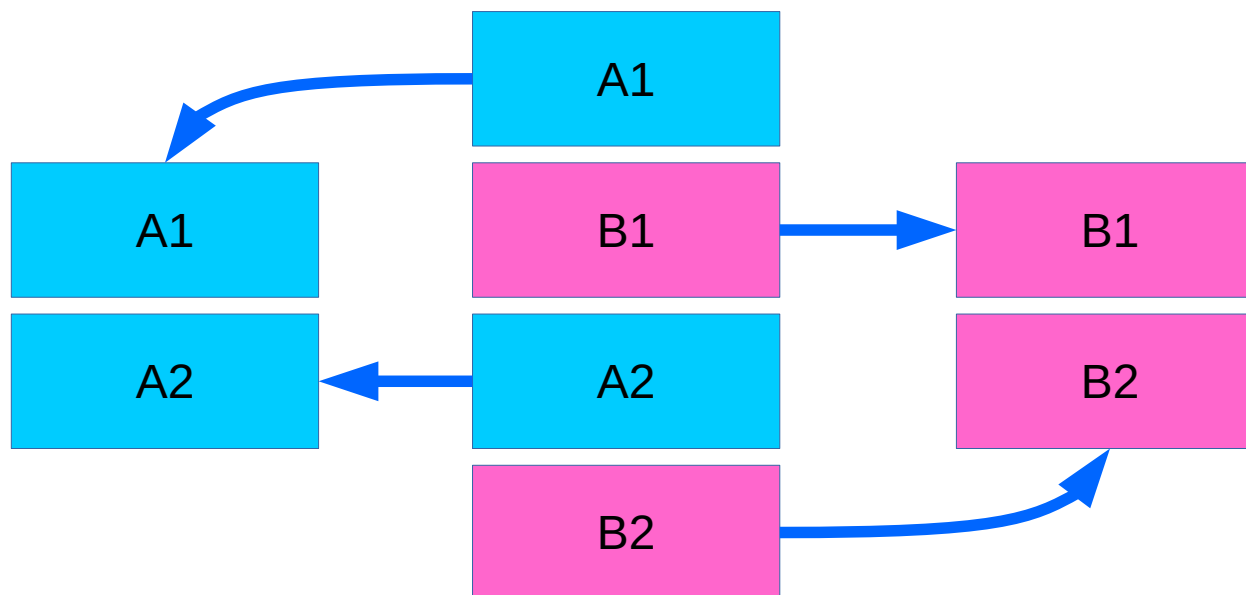
```
# A tibble: 6 x 4
```

	country	year	cases	population
	<chr>	<int>	<int>	<int>
1	Afghanistan	1999	745	19987071
2	Afghanistan	2000	2666	20595360
3	Brazil	1999	37737	172006362
4	Brazil	2000	80488	174504898
5	China	1999	212258	1272915272
6	China	2000	213766	1280428583



# separate(): **table3**

- What do I do if I know that my column contains mixed data entries?
- Given a regular expression for a delimiter, *separate()* turns a single character column into multiple columns.





# separate(): **table3**

```
table3 %>%  
  separate(rate,  
    into = c("cases",  
             "population"),  
    sep = "/" )
```

Here's how:  
**Push** the data  
*into* two columns

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

table3

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583



## Ex: Separating Compounded Entries: **table3**

	country	year	rate
1	Afghanistan	1999	745/19987071
2	Afghanistan	2000	2666/20595360
3	Brazil	1999	37737/172006362
4	Brazil	2000	80488/174504898
5	China	1999	212258/1272915272
6	China	2000	213766/1280428583

Break the string  
into length 2 chunks  
and place left in  
new col “*century*”  
and other in  
col “*year.*”

```
table3 %>%  
  separate(year, into = c("century", "year"), sep = 2)
```

```
table3 %>%  
  separate(rate, into = c("cases", "pop"), sep = "/")
```



# unite(): **table6**

```
table5 %>% unite(year,  
century, year, sep = "")
```

Here's how:  
**Pull** the data  
*from* two columns

country	<b>year</b>	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

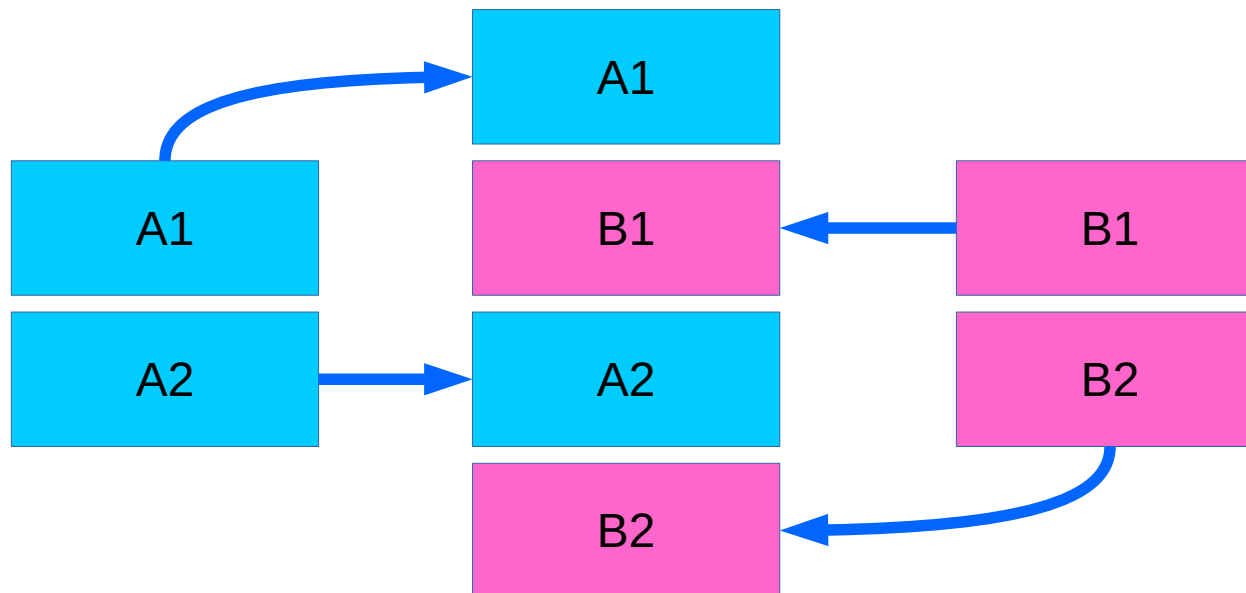
country	century	year	rate
Afghanistan	19	<b>99</b>	745 / 19987071
Afghanistan	20	<b>0</b>	2666 / 20595360
Brazil	19	<b>99</b>	37737 / 172006362
Brazil	20	<b>0</b>	80488 / 174504898
China	19	<b>99</b>	212258 / 1272915272
China	20	<b>0</b>	213766 / 1280428583

table6



## unite(): **table3**

- What do I do if I know that two columns contains data that could go into one column?
- Given a regular expression for pattern in text, `separate()` turns a single character column into multiple columns.







## Ex: Unite Compounded Entries

	country	century	year	rate
1	Afghanistan	19	99	745/19987071
2	Afghanistan	20	00	2666/20595360
3	Brazil	19	99	37737/172006362
4	Brazil	20	00	80488/174504898
5	China	19	99	212258/1272915272
6	China	20	00	213766/1280428583

What is the  
output of this?!

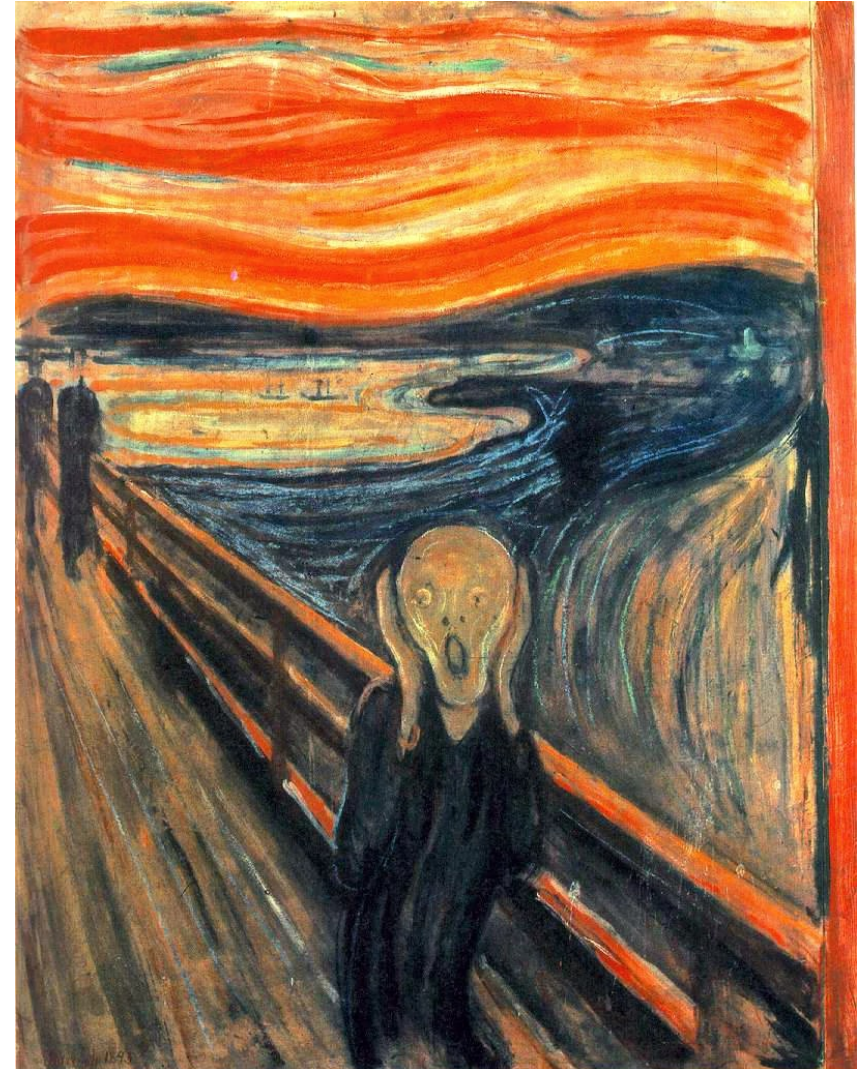
```
table5 %>%  
  unite(centuryYear,  
        century, year, sep = "")
```

# Missing Values!?

We may miss table entries

Two types of missing entries

- **Explicitly**, i.e., flagged with **NA**.
- **Implicitly**, i.e., simply not present in the data.



# Missing Data Illustrated With `tibble()`

# Make a table with a missing entry (NA).

```
stocks <- tibble(
```

```
  year = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),
```

```
  qtr  = c( 1, 2, 3, 4, 2, 3, 4),
```

```
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66))
```

Missing qtr:  
"1" for 2016

- Two missing values in this dataset:
  - The return for the **fourth** quarter of 2015 is explicitly missing, there is an entry of NA
  - The return for the first quarter of 2016 is implicitly missing, because it simply does not appear in the dataset.
- **Note: Missing data is easier to spot when viewing a table.**



# Missing Data In Table

	year	qtr	return
1	2015	1	1.88
2	2015	2	0.59
3	2015	3	0.35
4	2015	4	NA
5	2016	2	0.92
6	2016	3	0.17
7	2016	4	2.66

Missing "1"  
(first quarter)

Missing element

```
# Make a table with a missing entry (NA).
```

```
stocks <- tibble(
```

```
  year    = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),
```

```
  qtr     = c( 1,    2,    3,    4,    2,    3,    4),
```

```
  return  = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66))
```

# Spread the Missing Data

	qtr		2015	2016
1	1		1.88	NA
2	2		0.59	0.92
3	3		0.35	0.17
4	4		NA	2.66

Add NA elements  
to data set

```
# Make the implicit missing values explicit (i.e., adding  
NA's to make it clear that there is missing data).
```

```
# Use spread() to place both years into own column.
```

```
stocks %>%
```

```
  spread(year, return)
```



# Removing Missing Entries

```
# Remove all rows having "holes" in the data
```

```
# Create two cols for years 2015 and 2016
```

```
# Place years back into the same col again,  
removing the missing entries.
```

```
stocks %>%
```

```
spread(year, return) %>% gather(year, return,  
`2015`:`2016`, na.rm = TRUE)
```

```
> stocks %>%  
+   spread(year, return) %>% gather(year,  
return, `2015`:`2016`, na.rm = TRUE)  
# A tibble: 6 x 3  
  qtr  year return  
* <dbl> <chr> <dbl>  
1     1  2015   1.88  
2     2  2015   0.59  
3     3  2015   0.35  
4     2  2016   0.92  
5     3  2016   0.17  
6     4  2016   2.66
```

Are you throwing  
away your data?



The progression of the tables as the missing values are removed

Stocks

	year	qtr	return
1	2015	1	1.88
2	2015	2	0.59
3	2015	3	0.35
4	2015	4	NA
5	2016	2	0.92
6	2016	3	0.17
7	2016	4	2.66

1

Remove holes in rows

	qtr	year	return
1	1	2015	1.88
2	2	2015	0.59
3	3	2015	0.35
6	2	2016	0.92
7	3	2016	0.17
8	4	2016	2.66

3

Add NA

	qtr	2015	2016
1	1	1.88	NA
2	2	0.59	0.92
3	3	0.35	0.17
4	4	NA	2.66

2




## Let's Just Guess About The Missing Stuff... With `tribble()`

```
library(tibble)
#Create a table with missing entries
treatment <- tribble(
  ~ person, ~ treatment, ~response,
  "Derrick Whitmore", 1, 7,
  NA, 2, 10,
  NA, 3, 9,
  "Katherine Burke", 1, 4)
```

# Treatments Table With Missing Entries

- We assume that Derrick Whitmore's name makes up the missing entries.



	person	treatment	response
1	Derrick Whitmore	1	7
2	NA	2	10
3	NA	3	9
4	Katherine Burke	1	4



# Whitmore To The Rescue?

	person	treatment	response
1	Derrick Whitmore	1	7
2	Derrick Whitmore	2	10
3	Derrick Whitmore	3	9
4	Katherine Burke	1	4

**Can anything  
go wrong  
with this  
solution?!**

```
treatment %>%  
  fill(person)
```